

## Classifying Patronizing Comments in Stack Overflow

- Learning should be fun! Stack Overflow is a great way to have your questions answered by a community of fellow passionate scientists, researchers, or teachers. However, as many know, explaining things to people online can be very, very infuriating. Patronizing/condescending comments from frustrated users are common, and it makes for an unproductive and unfriendly experience. Hence, I want to use NLP to flag patronizing comments before they are sent, and send the user a request to choose their words again.
- Question: Is this a patronizing comment?
- Who benefits: stack overflow community

## Data Description:

- What dataset(s) do you plan to use, and how will you obtain the data?
  - Corpus (75.4m documents, each row is a comment): <https://www.kaggle.com/datasets/stackoverflow/stackoverflow?select=comments>
  - Condescending language research paper (essentially same project idea, but performed on reddit): <http://aclanthology.lst.uni-saarland.de/D19-1385.pdf>
    - Github w full dataset: <https://github.com/zijwang/talkdown>

## Tools:

- How do you intend to meet the tools requirement of the project?
  - Sklearn, nltk, vader, gensim, spacy, numpy, matplotlib, etc
  - Seaborn for visualization, maybe tableau if time allows

## MVP Goal:

- What would a minimum viable product (MVP) look like for this project?
  - A model to classify general negative replies, not necessarily condescending, which is more specific